



Writing Cache Friendly Code

——Slides derived from those by Randy Bryant ...





问题1

- Example: **cold cache**, 4-byte words, 4-word cache blocks

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```

Miss rate = ?

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

Miss rate = ?

Layout of C Arrays in Memory

C arrays allocated in **row-major** order


- each row in contiguous memory locations

which row
which column
 $A[i][j]$

$A[1][5]$

1 →

5 ↓



$A[1][5]$





问题2

- 一个 $n \times n$ 的二维数组，按行、列、对角线和反对角线访问，并且在不同的变址位移量情况下，都能实现无冲突访问？

A 4x4 array diagram illustrating memory access patterns. The array is represented as a grid of cells. The first column is labeled with a_{00} , a_{10} , a_{20} , and a_{30} from top to bottom, with a downward arrow to its left. The first row is labeled with a_{00} , a_{01} , a_{02} , and a_{03} from left to right, with a rightward arrow above it. A main diagonal line runs from the top-left cell (a_{00}) to the bottom-right cell (a_{33}). An anti-diagonal line runs from the top-right cell (a_{03}) to the bottom-left cell (a_{30}). The cells are labeled as follows:

a_{00}	a_{01}	a_{02}	a_{03}
a_{10}	a_{11}	a_{12}	a_{13}
a_{20}	a_{21}	a_{22}	a_{23}
a_{30}	a_{31}	a_{32}	a_{33}



Cache 性能

$$CPU\ time = \frac{IC \times \left(CPI_{Exec} + \frac{MemAccess}{Inst} \times MissRate \times MissPenalty \right)}{Clock\ rate}$$

$$CPU\ time = \frac{IC \times \left(\frac{ALUops}{Inst} \times CPI_{ALUops} + \frac{MemAccess}{Inst} \times AMAT \right)}{Clock\ rate}$$

$$\begin{aligned} AMAT &= Hit\ time + Miss\ Rate \times Miss\ Penalty \\ &= \% instructions \times (Hit\ time_{Inst} + Miss\ Rate_{inst} \times Miss\ Penalty_{Inst}) \\ &\quad + \% data \times (Hit\ time_{Data} + Miss\ Rate_{Data} \times Miss\ Penalty_{Data}) \end{aligned}$$



Cache性能

■ Cache 优化

- 1. 减少缺失率 (miss rate)
- 2. 减少缺失代价 (miss penalty)
- 3. 较少判断命中时间 (hit time)

$$AMAT = HitTime + MissRate \times MissPenalty$$



Cache Organization?

- Change **Block Size**
- Change **Cache Size**
- Change **Cache Internal Organization**
- Change **Associativity**
- Change **Compiler**
- Which of 3Cs is obviously affected?



Cache的设计空间

- Larger cache size
 - + reduces capacity and conflict misses
 - - hit time will increase
- Higher associativity
 - + reduces conflict misses
 - - may increase hit time
- Larger block size
 - + reduces compulsory and capacity misses
 - + exploit burst transfers in memory and on buses
 - - increases conflict misses and miss penalty



减少缺失率(MissRate)

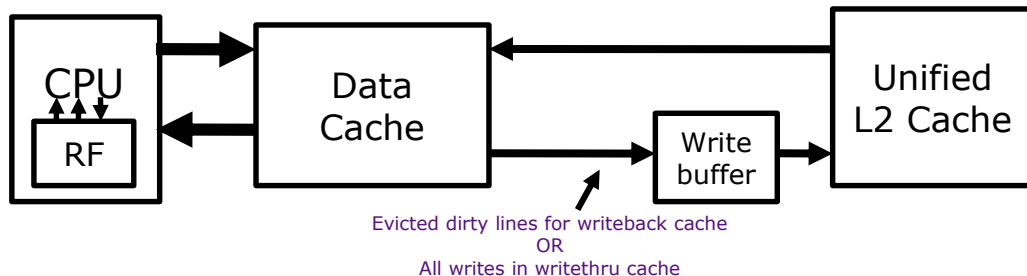
$$CPU\ time = \frac{IC \times \left(CPI_{Exec} + \frac{MemAccess}{Inst} \times MissRate \times MissPenalty \right)}{Clock\ rate}$$

1. Larger Cache => Reduce Capacity
2. Larger Block Size => Reduce Compulsory
3. Higher Associativity => Reduce Confilcts
4. Way Prediction & Pseudo-Associativity
5. Compiler Optimizations



减少缺失代价(MissPenalty)


- 1. Multilevel Caches
- 2. Critical Word First and Early Restart
- 3. Giving Priority to Read Misses over Writes
- 4. Merging Write Buffer
- 5. Victim Caches
- * Reducing Read Miss Penalty—Write Buffer



$$CPU \text{ time} = \frac{IC \times \left(CPI_{Exec} + \frac{MemAccess}{Inst} \times MissRate \times MissPenalty \right)}{Clock \text{ rate}}$$



多级Cache



- **Inclusive multilevel cache:**

- Inner cache holds copies of data in outer cache
 - External access need only check outer cache
 - Most common case

- **Exclusive multilevel caches:**


- Inner cache may hold data not in outer cache
 - Swap lines between inner/outer caches on miss
 - Used in AMD Athlon with 64KB primary and 256KB secondary cache



减少命中时间 (HitTime)

- Small and Simple Caches
- Avoiding Address Translation
- Reducing Write Hit Time—Pipelining Cache Write

$$AMAT = \text{HitTime} + \text{MissRate} \cdot \text{MissPenalty}$$

	Technique	Miss-Rate	Miss-Penalty	HitTime	Complexity
	Larger Block Size	+	-		0
	Higher Associativity	+		-	1
	Victim Caches	+			2
	Pseudo-Associative Caches	+			2
	HW Prefetching of Instr/Data	+	+		2
	Compiler Controlled Prefetching	+	+		3
	Compiler Reduce Misses				0
	Priority to Read Misses		+		1
	Early Restart & Critical Word 1st		+		2
	Non-Blocking Caches		+		3
	Second Level Caches		+		2
	Better memory system		+		3
	Small & Simple Caches	-		+	0
	Avoiding Address Translation			+	2
Computer Arch	Pipelining Caches			+	2





Instruction vs. Data Caches

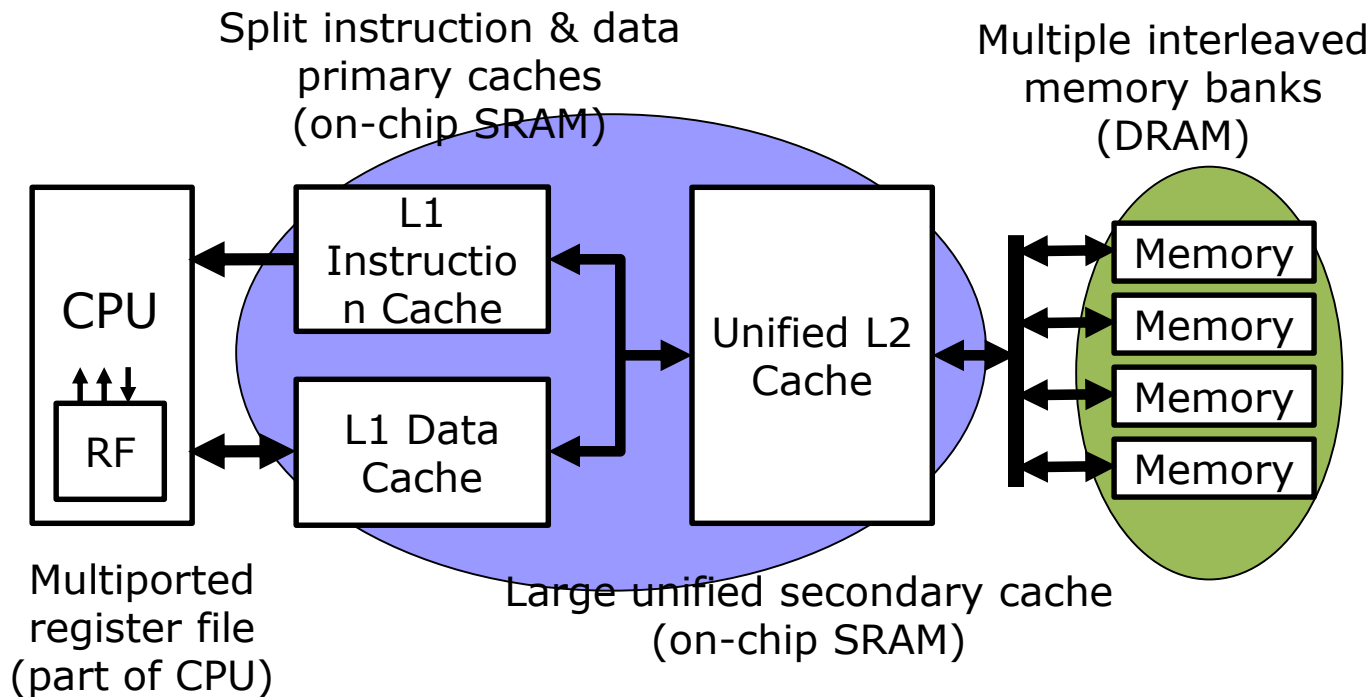
- **Separate or Unified?**
- **Unified:**
 - + Dynamic sharing of cache space: no over provisioning that might happen with static partitioning (i.e., split I and D caches)
 - Instructions and data can thrash each other (i.e., no guaranteed space for either)
 - I and D are accessed in different places in the pipeline. Where do we place the unified cache for fast access?
- **First level caches are almost always split**
 - Mainly for the last reason above
- **Second and higher levels are almost always unified**



Multi-level Caching in a Pipelined Design

- **First-level caches (instruction and data)**
 - Decisions very much affected by cycle time
 - Small, lower associativity
 - **Tag store and data store accessed in parallel**
- **Second-level caches**
 - Decisions need to balance hit rate and access latency
 - Usually large and highly associative; **latency not as important**
 - **Tag store and data store accessed serially**
- **Serial vs. Parallel access of levels**
 - **Serial: Second level cache accessed only if first-level misses**
 - **Second level does not see the same accesses as the first**
 - First level acts as a filter (filters some temporal and spatial locality)
 - Management policies are therefore different

存储器系统层次结构





Performance of 1-Level Split Cache

$$\begin{aligned}CPI^{stall} &= \frac{\text{instruction stall cycles}}{\text{stalls}} \times \frac{\text{instruction stalls}}{\text{instruction access}} \times \frac{\text{instruction accesses}}{\text{instruction}} \\&\quad + \frac{\text{data stall cycles}}{\text{stalls}} \times \frac{\text{data stalls}}{\text{data access}} \times \frac{\text{data accesses}}{\text{instruction}} \\&= \text{instruction miss penalty} \times \text{instruction miss rate} \times \frac{1 \text{ instruction access}}{\text{instruction}} \\&\quad + \text{data miss penalty} \times \text{data miss rate} \times \frac{IC^{load} + IC^{store}}{IC}\end{aligned}$$

Typical values

Instruction Miss Penalty = Data Miss Penalty = 50 cycles per stall

Instruction Miss Rate = 0.5%

Data Miss Rate = 5.0%

$IC^{load} = 25\%$



$IC^{store} = 15\%$

$$CPI^{stall} = 50 \times 0.005 \times 1 + 50 \times 0.05 \times 0.40 = 0.25$$

$CPI = 1.25 \Rightarrow 20\%$ degradation



Performance of Two-Level Unified Cache




$$CPI^{\text{stall}} = \frac{\text{stall cycles}}{IC} = \sum_{n = \text{stall types}} \frac{\text{stall cycles}}{\text{stall type } n} \times \frac{\text{stalls of type } n}{IC}$$

$$\text{stall types} = \sum_{i = (\text{instruction, data})} \{ (L1 \text{ miss, L2 hit})_i, (L1 \text{ miss, L2 miss})_i \}$$



Stalls in Two-Level Unified Cache



$$\begin{aligned} \text{CPI}^{\text{stall}} &= \sum_{i = \{\text{data}, \text{instr}\}} \left[\frac{\text{stall cycles}}{(\text{L1 miss, L2 hit})_i} \times \frac{(\text{L1 miss, L2 hit})_i}{\text{IC}} \right. \\ &\quad \left. + \frac{\text{stall cycles}}{(\text{L1 miss, L2 miss})_i} \times \frac{(\text{L1 miss, L2 miss})_i}{\text{IC}} \right] \\ &= \sum_{i = \{\text{data}, \text{instr}\}} \left[\frac{\text{stall cycles}}{(\text{L1 miss, L2 hit})_i} \times \frac{(\text{L1 miss, L2 hit})_i}{\text{L1 access}_i} \times \frac{\text{L1 access}_i}{\text{IC}} \right. \\ &\quad \left. + \frac{\text{stall cycles}}{(\text{L1 miss, L2 miss})_i} \times \frac{(\text{L1 miss, L2 miss})_i}{\text{L1 access}_i} \times \frac{\text{L1 access}_i}{\text{IC}} \right] \end{aligned}$$



Assume L1 and L2 Are Statistically Independent

$$\begin{aligned}\frac{(L1 \text{ miss, L2 hit})_i}{L1 \text{ access}_i} &= \frac{L1 \text{ miss}_i \text{ followed by L2 hit}}{L1 \text{ access}_i} = \frac{L1 \text{ miss}_i}{L1 \text{ access}_i} \times \underbrace{\frac{L2 \text{ access}}{L1 \text{ miss}_i}}_1 \times \frac{L2 \text{ hit}}{L2 \text{ access}} \\ &= \frac{L1 \text{ miss}_i}{L1 \text{ access}_i} \times \frac{L2 \text{ hit}}{L2 \text{ access}}\end{aligned}$$

$$\begin{aligned}\frac{(L1 \text{ miss, L2 miss})_i}{L1 \text{ access}_i} &= \frac{L1 \text{ miss}_i \text{ followed by L2 miss}}{L1 \text{ access}_i} = \frac{L1 \text{ miss}_i}{L1 \text{ access}_i} \times \underbrace{\frac{L2 \text{ access}}{L1 \text{ miss}_i}}_1 \times \frac{L2 \text{ miss}}{L2 \text{ access}} \\ &= \frac{L1 \text{ miss}_i}{L1 \text{ access}_i} \times \frac{L2 \text{ miss}}{L2 \text{ access}}\end{aligned}$$

$$\begin{aligned}CPI^{\text{stall}} = \sum_{i = \{\text{data, instr}\}} &\left[\frac{\text{stall cycles}}{(L1 \text{ miss, L2 hit})_i} \times \frac{L1 \text{ miss}_i}{L1 \text{ access}_i} \times \frac{L2 \text{ hit}}{L2 \text{ access}} \times \frac{L1 \text{ access}_i}{IC} \right. \\ &\left. + \frac{\text{stall cycles}}{(L1 \text{ miss, L2 miss})_i} \times \frac{L1 \text{ miss}_i}{L1 \text{ access}_i} \times \frac{L2 \text{ miss}}{L2 \text{ access}} \times \frac{L1 \text{ access}_i}{IC} \right]\end{aligned}$$



Definitions


Miss Rates

$$\begin{aligned}M_{L1} &= \text{Miss Rate at L1} \\ &= \frac{\text{L1 miss}}{\text{L1 access}}\end{aligned}$$

$$\begin{aligned}1 - M_{L1} &= \text{Hit Rate at L1} \\ &= \frac{\text{L1 hit}}{\text{L1 access}}\end{aligned}$$

$$\begin{aligned}M_{L2} &= \text{Miss Rate at L2} \\ &= \frac{\text{L2 miss}}{\text{L1 access}}\end{aligned}$$

$$\begin{aligned}1 - M_{L2} &= \text{Hit Rate at L2} \\ &= \frac{\text{L2 hit}}{\text{L2 access}}\end{aligned}$$



Miss Penalties

$$\begin{aligned}P_{L1} &= \text{Miss Penalty (cycles) from L1 to L2} \\ &= \frac{\text{stall cycles}}{(\text{L1 miss, L2 hit})}\end{aligned}$$

$$P_{L2} = \text{Miss Penalty (cycles) from L2 to Main Memory}$$

$$P_{L1} + P_{L2} = \frac{\text{stall cycles}}{(\text{L1 miss, L2 miss})}$$

IC^A = Data Access Instructions (Load/Store)

IC = Total Instructions



Simplifying

$$\begin{aligned} \text{CPI}^{\text{stall}} &= \sum_{i = \{\text{data}, \text{instr}\}} \left[\frac{\text{stall cycles}}{(\text{L1 miss, L2 hit})_i} \times \frac{\text{L1 miss}_i}{\text{L1 access}_i} \times \frac{\text{L2 hit}}{\text{L2 access}} \times \frac{\text{L1 access}_i}{\text{IC}} \right. \\ &\quad \left. + \frac{\text{stall cycles}}{(\text{L1 miss, L2 miss})_i} \times \frac{\text{L1 miss}_i}{\text{L1 access}_i} \times \frac{\text{L2 miss}}{\text{L2 access}} \times \frac{\text{L1 access}_i}{\text{IC}} \right] \\ &= P_{L1} \times M_{L1} \times (1 - M_{L2}) \times \frac{\text{IC} + \text{IC}^A}{\text{IC}} + (P_{L1} + P_{L2}) \times M_{L1} \times M_{L2} \times \frac{\text{IC} + \text{IC}^A}{\text{IC}} \\ &= M_{L1} \times \left[1 + \frac{\text{IC}^A}{\text{IC}} \right] \times [P_{L1} \times (1 - M_{L2}) + (P_{L1} + P_{L2}) \times M_{L2}] \\ &= M_{L1} \times \left[1 + \frac{\text{IC}^A}{\text{IC}} \right] \times [P_{L1} - P_{L1} \times M_{L2} + P_{L1} \times M_{L2} + P_{L2} \times M_{L2}] \\ \text{CPI}^{\text{stall}} &= M_{L1} \times \left[1 + \frac{\text{IC}^A}{\text{IC}} \right] \times [P_{L1} + P_{L2} \times M_{L2}] \end{aligned}$$



Splitting L1 Cache with Unified L2 Cache

$$\begin{aligned}\frac{\text{misses at L1}}{IC} &= \frac{\text{misses at L1}}{\text{accesses at L1}} \times \frac{\text{accesses at L1}}{IC} \\ &= M_{L1} \times \left[1 + \frac{IC^A}{IC} \right] = M_{L1} + M_{L1} \times \frac{IC^A}{IC}\end{aligned}$$

↓

$$\begin{aligned}\frac{\text{data misses at L1}}{\text{data accesses at L1}} + \frac{\text{instruction misses at L1}}{\text{instruction accesses at L1}} &= \sum_{i \in (\text{data}, \text{instructions})} \frac{\text{misses at L1}_i}{\text{accesses at L1}_i} \times \frac{\text{accesses at L1}_i}{IC} \\ &= M_{L1}^{\text{instructions}} \times 1 + M_{L1}^{\text{data}} \times \frac{IC^A}{IC} \\ &= M_{L1}^I + M_{L1}^D \times \frac{IC^A}{IC}\end{aligned}$$

$$CPI_{\text{unified}}^{\text{stall}} = M_{L1} \times \left[1 + \frac{IC^A}{IC} \right] \times [P_{L1} + P_{L2} \times M_{L2}]$$

↓

$$CPI_{\text{split}}^{\text{stall}} = \left[M_{L1}^I + M_{L1}^D \times \frac{IC^A}{IC} \right] \times [P_{L1} + P_{L2} \times M_{L2}]$$



Second Layer Cache (L2) with Split L1

Definitions

P_{L1} = Miss Penalty (cycles) at L1

P_{L2} = Miss Penalty (cycles) at L2

IC^A = Data Access Instructions (Load/Store)

IC = Total Instructions

M_{L1}^I = Instruction Miss Rate at L1

M_{L1}^D = Data Miss Rate at L1

M_{L2} = Miss Rate at L2

One layer (L1) of split cache

$$CPI_{1-level}^{stall} = \left[M_{L1}^I + M_{L1}^D \times \frac{IC^A}{IC} \right] \times P_{L1}$$


Miss Penalty at L1 (to main memory) $P_{L1} \sim 50$ cycles

Split L1 cache and unified L2 cache

Miss Penalty at L1 (to L2) $P_{L1} \sim 5$ cycles

Miss Penalty at L2 (to main memory) $P_{L2} \sim 50$ cycles

Miss Rate at L2 $\sim 1\%$



$$CPI_{2-level}^{stall} = \left[M_{L1}^I + M_{L1}^D \times \frac{IC^A}{IC} \right] \times (P_{L1} + P_{L2} \times M_{L2})$$

$$P_{L1} + P_{L2} \times M_{L2} = 5 + 50 \times 0.01 = 5.5$$



Miss Example — Extreme Locality

Program

```
int i, a;  
for (i = 0 ; i < 4096 ; i++){  
    a = i;  
}
```

Compiler assignments

Register $\leftarrow i$

Memory $\leftarrow a$

Memory accesses

4096 write accesses to integer **a**

1 compulsory cache miss (write allocate) on $i = 0$

4095 cache hits on $i > 0$

0 read accesses to **a**

Miss rate

$$\text{miss rate} = \frac{1 \text{ miss}}{4096 \text{ accesses}} = 2.44 \times 10^{-4}$$



Miss Example — Extreme Non-Locality

Block size

16 bytes/block = 4 words/block = 4 array elements

每一块可以存放数组的4个元素

Program

```
int i, a[4096];  
for (i = 0 ; i < 4096 ; i++){  
    a[i] = i;  
}
```

Compiler assignments

Register \leftarrow i
Memory \leftarrow a[]

Memory accesses

4096 write accesses to integer array a[]

Compulsory cache miss every 4 array elements

4096/4 = 1024 cache misses (write allocate)

0 read accesses to a[]

Miss rate

$$\text{miss rate} = \frac{1024 \text{ misses}}{4096 \text{ accesses}} = 0.25$$

a[i][j]	j=0	j=1	j=2	j=3	j=4	j=5	j=6	j=7
i=0	1 [m]	2 [h]	3 [h]	4 [h]	5 [m]	6 [h]	7 [h]	8 [h]
i=1	9 [m]	10 [h]	11 [h]	12 [h]	13 [m]	14 [h]	15 [h]	16 [h]
i=2	17 [m]	18 [h]	19 [h]	20 [h]	21 [m]	22 [h]	23 [h]	24 [h]
i=3	25 [m]	26 [h]	27 [h]	28 [h]	29 [m]	30 [h]	31 [h]	32 [h]





Miss Example — Good Locality

Cache parameters	B	S	W	Capacity = B × S × W
	16	256	4	16 KB = 4 Kwords

4路组相连
每组256块
每块4个字 (数组元素) (16字节)

Program

```
int a[4096], i, j;  
for (i = 0 ; i < 10 ; i++){  
    for (j = 0 ; j < 4096 ; j++){  
        a[j] = i + j;  
    }  
}
```

Memory accesses

40960 write accesses to integer array **a[]**

$i = 0$: Compulsory cache miss every 4 elements

$i > 0$: Entire array in cache \Rightarrow cache hits

Miss rate

$$\text{miss rate} = \frac{1024 \text{ misses}}{40960 \text{ accesses}} = 0.025$$

Compiler assignments

Register $\leftarrow i, j$

Memory $\leftarrow a[]$

	255	511	767	1023

↑ sets	2	259	514	770
	1	257	513	769
	0	256	512	768
	0	1	2	3
	4 slots			

4K integers = 16 Kbytes
= 1024 blocks



Miss Example — Better Locality Using MRU (Most Recently Used)

Cache parameters	B	S	W	Capacity = B × S × W
	16	256	4	16 KB = 4 Kwords

4路组相连
每组256块
每块4个字 (数组元素) (16字节)

Program

```
int a[5120], i, j;
for (i = 0 ; i < 10 ; i++) {
    for (j = 0 ; j < 5120 ; j++) {
        a[j] = i + j;
    }
}
```

Compiler assignments

Register ← i, j
Memory ← a[]

255	511	767	1023	
...	
2	259	514	770	
1	257	513	769	
0	256	512	768	
0	1	2		3

4 slots

5K integers = 1280 blocks

Memory accesses with MRU

51200 write accesses to integer array **a[]**

i = 0: Compulsory cache miss every 4 elements

i > 0: Conflict misses to slot 3 (2 out of 5 accesses)

j=0...255, 1024...1279

Miss rate

$$\text{miss rate} = \frac{1280 + 512 \times 9 \text{ misses}}{51200 \text{ accesses}} = 0.025 \times \left(1 + \frac{2}{5} \times 9 \right) = 0.115$$



Miss Example — Improved Locality

Cache parameters	B	S	W	Capacity = B × S × W
	16	256	4	16 KB = 4 Kwords

Program

```
int a[5120], i, j;  
for (i = 0 ; i < 10 ; i++){  
    for (j = 0 ; j < 4096 ; j++){  
        a[j] = i + j;  
    }  
    for (i = 0 ; i < 10 ; i++){  
        for (j = 4096 ; j < 5120 ; j++){  
            a[j] = i + j;  
        }  
    }  
}
```

Memory accesses

51200 write accesses to integer array **a[]**

i = 0: Compulsory cache miss every 4 elements

i > 0: Entire array in cache ⇒ cache hits

Miss rate

$$\text{miss rate} = \frac{1024 + 256 \text{ misses}}{40960 + 10240 \text{ accesses}} = 0.025$$

Compiler assignments

Register ← **i, j**

Memory ← **a[]**

255	511	767	1023
...
2	259	514	770
1	257	513	769
0	256	512	768

1279	511	767	1023
...
1026	259	514	770
1025	257	513	769
1024	256	512	768

5K integers = 1280 blocks



Miss Example — Address Aliasing

Cache parameters	B	S	W	Capacity = B × S × W
	16	256	2	8 KB = 2 Kwords

Program

```
int a[512], b[512], c[512], i, j;
for (i = 0 ; i < 20; i++){
    for (j = 0 ; j < 512 ; j++){
        c[j] = a[j] + b[j] + c[j];
    }
}
```

Compiler assignments

Register $\leftarrow i, j$

Memory (address = $0200AS_1S_2B$)

a: $02000000 - 020007FF$

b: $02001000 - 020017FF$

c: $02002000 - 020027FF$

3×512 integers = 6 KB = 384 blocks

Memory accesses

$20 \times 3 \times 512 = 30720$ read accesses to $a[]$, $b[]$, $c[]$

$20 \times 512 = 10240$ write accesses to array $c[]$


Set assignment = $\text{int}(0200AS_1S_2B/10) \% 100 = S_1S_2$

$i = 0$: $3 \times 512 / 4 = 384$ compulsory misses

$i > 0$: $2 \times 512 / 4 = 256$ conflict misses for sets $a[]$, $c[]$

Miss rate

$$\text{miss rate} = \frac{384 + 256 \times 19 \text{ misses}}{30720 + 10240 \text{ accesses}} = \frac{2688}{40960} = 0.128$$





Miss Example — Address Aliasing with Larger W

Cache parameters	B	S	W	Capacity = B × S × W
	16	128	4	8 KB = 2 Kwords

Program

```
int a[512], b[512], c[512], i, j;
for (i = 0 ; i < 20; i++){
    for (j = 0 ; j < 512 ; j++){
        c[j] = a[j] + b[j] + c[j];
    }
}
```

Compiler assignments

Register ← i, j

Memory (address = 0200AS₁S₂B)

a: 02000000 – 020007FF

b: 02001000 – 020017FF

c: 02002000 – 020027FF

3 × 512 integers = 6 KB = 384 blocks

Memory accesses

20 × 3 × 512 = 30720 read accesses to a[], b[], c[]

20 × 512 = 10240 write accesses to array c[]


Set assignment = $\text{int}(0200AS_1S_2B/10) \% 80 = S_1S_2$

i = 0: 3 × 512 / 4 = 384 compulsory misses

i > 0: All arrays in cache ⇒ cache hits

Miss rate

$$\text{miss rate} = \frac{384 \text{ misses}}{30720 + 10240 \text{ accesses}} = \frac{384}{40960} = 0.009375$$





Writing Cache Friendly Code

Two major rules:

- Repeated references to data are good (**temporal locality**)
- Stride reference patterns are good (**spatial locality**)
- Example: **cold cache**, 4-byte words, 4-word cache blocks

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];
    return sum;
}
```

Miss rate = $1/4 = 25\%$

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;

    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];
    return sum;
}
```

Miss rate = 100%



Layout of C Arrays in Memory

C arrays allocated in **row-major** order


- each row in contiguous memory locations

which row
which column
 $A[i][j]$

$A[1][5]$

1 →

5 ↓





Layout of C Arrays in Memory

C arrays allocated in **row-major** order

- each row in contiguous memory locations

Stepping through columns in one row:

- `for (i = 0; i < N; i++)`
 `sum += a[0][i];`
- **accesses successive elements**
- **if block size (B) > 4 bytes, exploit spatial locality**
 - compulsory miss rate = 4 bytes / B

Stepping through rows in one column:

- `for (i = 0; i < n; i++)`
 `sum += a[i][0];`
- **accesses distant elements**
- **no spatial locality!**
 - compulsory miss rate = 1 (i.e. 100%)





Matrix Multiplication Example

Major Cache Effects to Consider

□ Total cache size

- Exploit temporal locality and keep the working set small (e.g., use blocking)

□ Block size

- Exploit spatial locality

Description:

□ Multiply $N \times N$ matrices

□ $O(N^3)$ total operations

□ Accesses

- N reads per source element
- N values summed per destination
 - but may be able to hold in register

```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

Variable *sum*
held in register

/




Miss Rate Analysis for Matrix Multiply

Assume:

- Block size = 32B (big enough for **four** 64-bit words)
- Matrix dimension (N) is very large
 - Approximate $1/N$ as 0.0
- Cache is not even big enough to hold multiple rows

Analysis Method:

- Look at access pattern of inner loop






Matrix Multiplication (ijk)

```
/* ijk */  
for (i=0; i<n; i++) {  
  for (j=0; j<n; j++) {  
    sum = 0.0;  
    for (k=0; k<n; k++)  
      sum += a[i][k] * b[k][j];  
    c[i][j] = sum;  
  }  
}
```

Inner loop:



Misses per Inner Loop Iteration:

A
0.25

B
1.0


C
0.0 (=1/N)



Matrix Multiplication (jik)

```
/* jik */  
for (j=0; j<n; j++) {  
  for (i=0; i<n; i++) {  
    sum = 0.0;  
    for (k=0; k<n; k++)  
      sum += a[i][k] * b[k][j];  
    c[i][j] = sum  
  }  
}
```

Inner loop:



Misses per Inner Loop Iteration:

A
0.25

B
1.0


C
0.0



Matrix Multiplication (kij)

```
/* kij */  
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

Inner loop:



Misses per Inner Loop Iteration:

A
0.0

B
0.25


C
0.25



Matrix Multiplication (ikj)

```
/* ikj */  
for (i=0; i<n; i++) {  
    for (k=0; k<n; k++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

Inner loop:



Misses per Inner Loop Iteration:

A
0.0


B
0.25

C
0.25

Matrix Multiplication (jki)


```
/* jki */
for (j=0; j<n; j++) {
    for (k=0; k<n; k++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
```

Inner loop:




A

Column -
wise



B

Fixed



C

Column-
wise

Misses per Inner Loop Iteration:

A
1.0


B
0.0

C
1.0



Matrix Multiplication (kji)

```
/* kji */  
for (k=0; k<n; k++) {  
    for (j=0; j<n; j++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```



Misses per Inner Loop Iteration:

A
1.0

B
0.0

C
1.0



Summary of Matrix Multiplication

```
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

ijk (& jik):

- 2 loads, 0 stores
- misses/iter = 1.25

```
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

kij (& ikj):

- 2 loads, 1 store
- misses/iter = 0.5

```
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

jki (& kji):

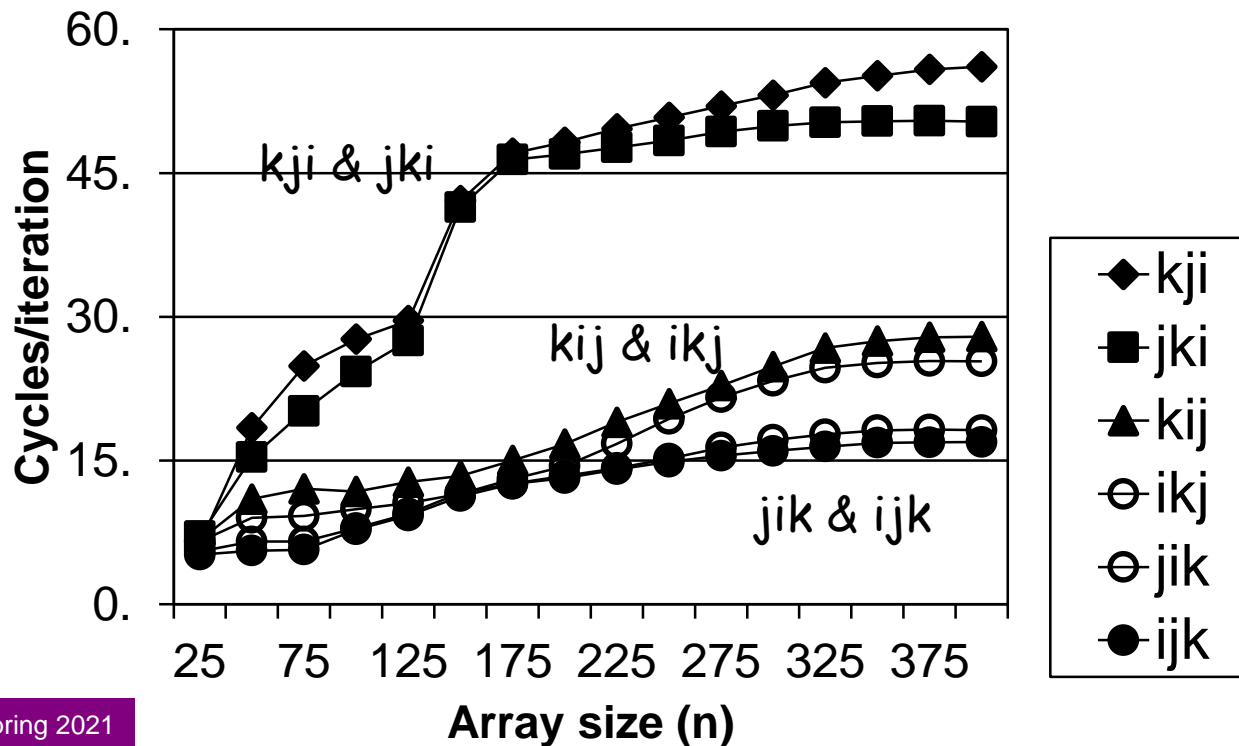
- 2 loads, 1 store
- misses/iter = 2.0



Pentium Matrix Multiply Performance

Miss rates are helpful but not perfect predictors.

- Code scheduling matters, too.





Improving Temporal Locality by Blocking

Example: Blocked matrix multiplication

- “**block**” (in this context) does not mean “**cache block**”.
- Instead, it mean a sub-block within the matrix.
- Example: $N = 8$; sub-block size = 4

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix}$$

Key idea: Sub-blocks (i.e., A_{xy}) can be treated just like scalars.

$$C_{11} = A_{11}B_{11} + A_{12}B_{21} \quad C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21} \quad C_{22} = A_{21}B_{12} + A_{22}B_{22}$$

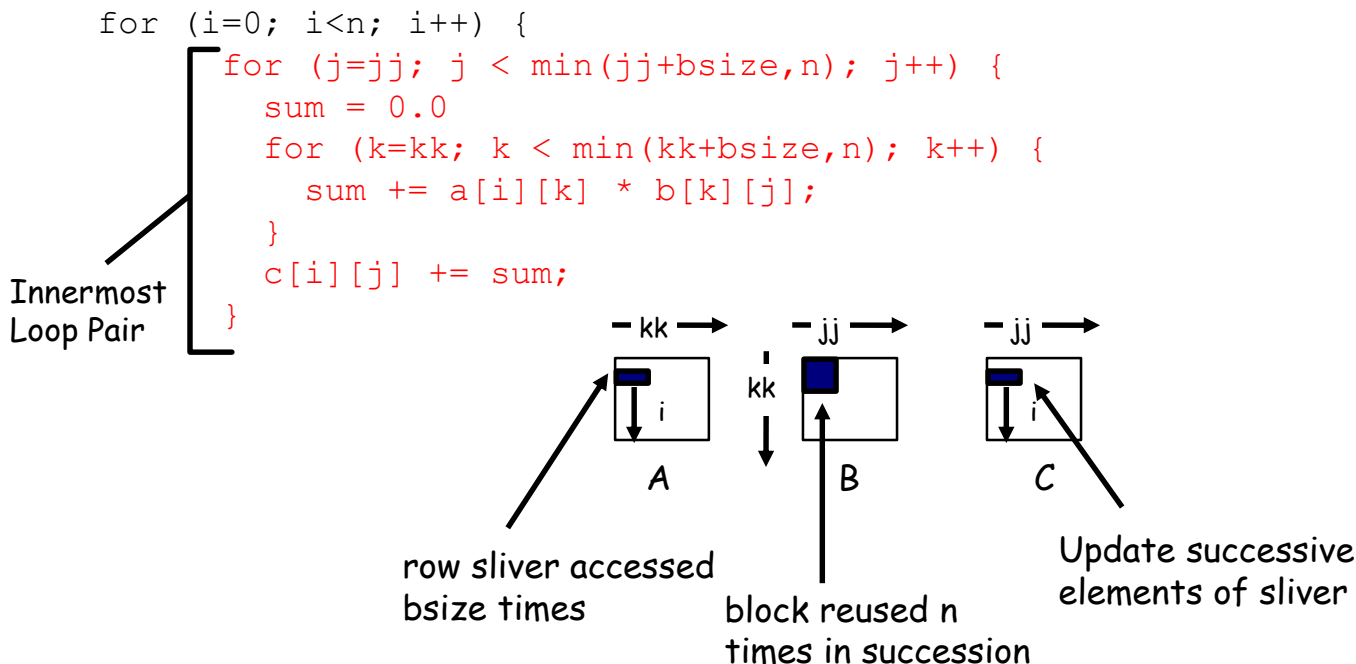


Blocked Matrix Multiply (bijk)

```
for (jj=0; jj<n; jj+=bsize) {  
  
    for (i=0; i<n; i++)  
        for (j=jj; j < min(jj+bsize,n); j++)  
            c[i][j] = 0.0;  
  
    for (kk=0; kk<n; kk+=bsize) {  
        for (i=0; i<n; i++) {  
            for (j=jj; j < min(jj+bsize,n); j++) {  
                sum = 0.0  
                for (k=kk; k < min(kk+bsize,n); k++) {  
                    sum += a[i][k] * b[k][j];  
                }  
                c[i][j] += sum;  
            }  
        }  
    }  
}
```

Blocked Matrix Multiply Analysis

- Innermost loop pair multiplies a $1 \times bsize$ sliver of A by a $bsize \times bsize$ block of B and accumulates into $1 \times bsize$ sliver of C
- Loop over i steps through n row slivers of A & C , using same B






Pentium Blocked Matrix Multiply Performance

Blocking (bijk and bikj) improves performance by a factor of two over unblocked versions (ijk and jik)

- relatively insensitive to array size.





Concluding Observations

Programmer can optimize for cache performance

- How **data structures** are organized
- How data are accessed
 - Nested loop structure
 - **Blocking** is a general technique

All systems favor “cache friendly code”

- Getting absolute optimum performance is very platform specific
 - Cache sizes, line sizes, associativities, etc.
- Can get most of the advantage with generic code
 - Keep working set reasonably small (temporal locality)
 - Use small strides (spatial locality)