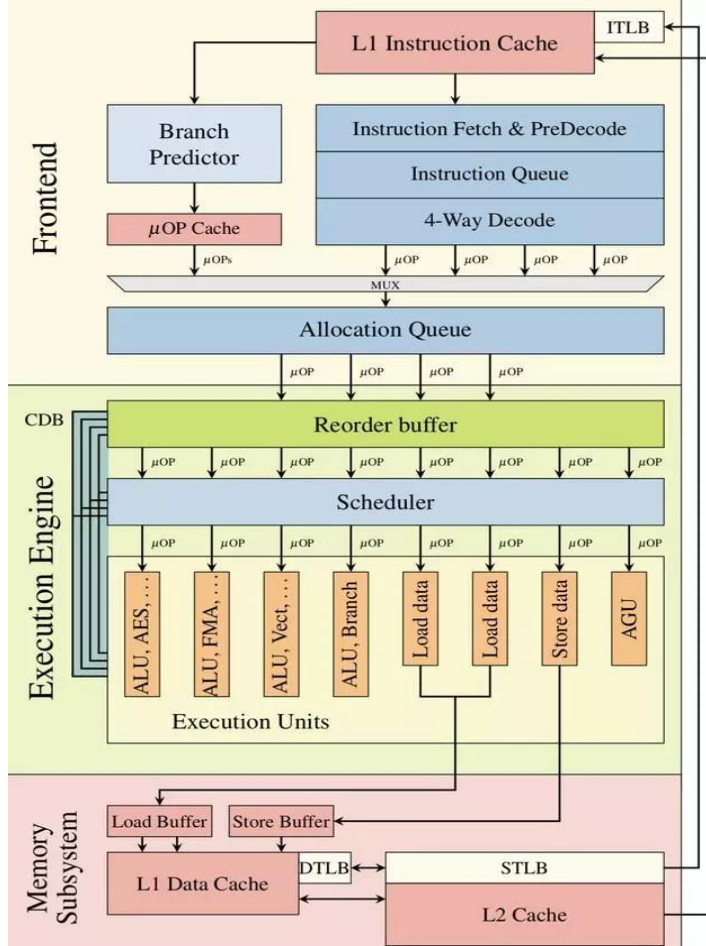
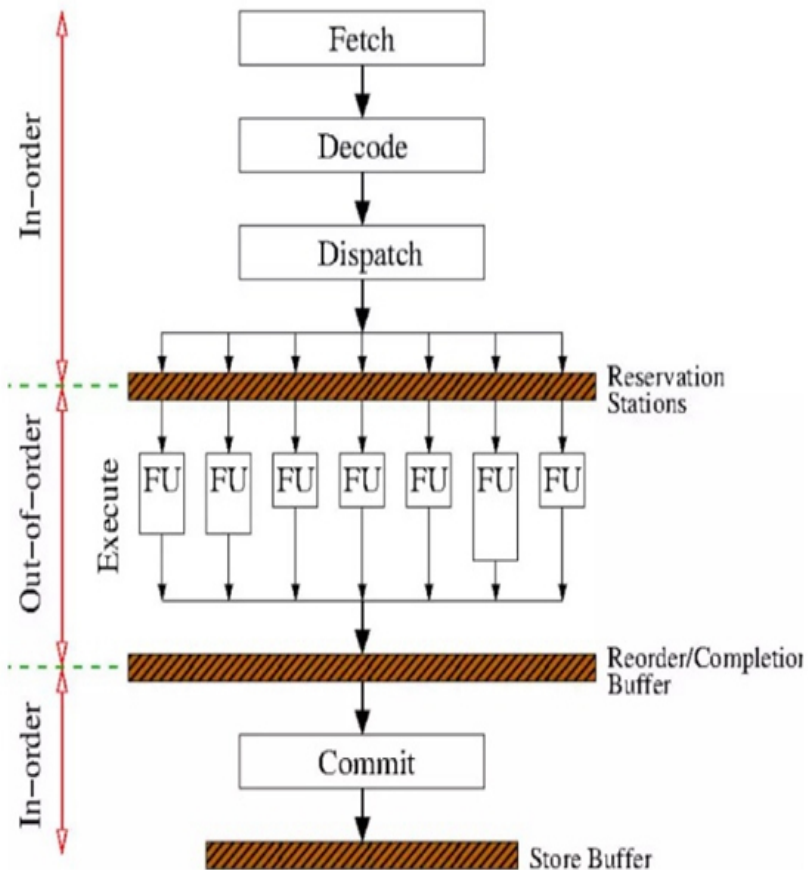




# Branch Prediction





熔断利用了图中execution engine的部分，幽灵利用了图中branch predictor的部分



# Branches

- Instructions which can alter the flow of instruction execution in a program



# Branch Prediction is **More** Important Today

**Conditional branches still comprise about 20% of instructions**

pipelines deeper branch not resolved until more cycles from fetching therefore the misprediction penalty greater

**cycle times smaller:** more emphasis on throughput (performance), more functionality between fetch & execute

multiple instruction issue (superscalars & VLIW) branch occurs almost every cycle flushing & refetching more instructions

**object-oriented programming** more indirect branches which harder to predict

dual of Amdahl's Law other forms of pipeline stalling are being addressed so the portion of CPI due to branch delays is relatively larger

On the other hand,

**chips are denser so we can consider sophisticated HW solutions**

**hardware cost is small compared to the performance gain**

**All this means that the potential stalling due to branches is greater**



# Types of Branches

	Conditional	Unconditional
Direct	if - then- else for loops (bez, bnez, etc)	procedure calls (jal) goto (j)
Indirect		return (jr) virtual function lookup function pointers (jalr)



问题1: 采用循环程序计算  $S=1+2+3+\dots+99+100$

采用1位预测器, 预测准确率是多少?

采用2位预测器, 预测准确率是多少?

■问题2: 饱和计数器 ( **saturating counter** ) ?

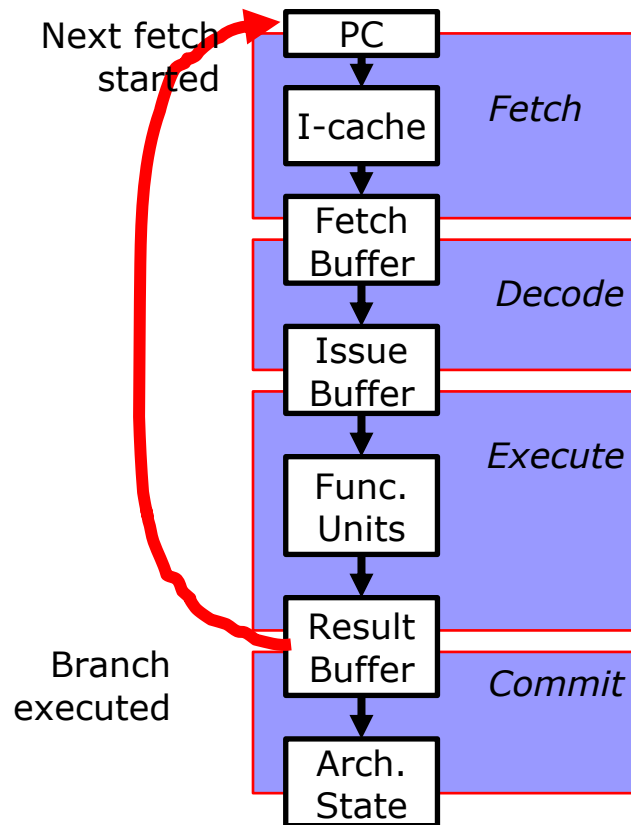


# Control Flow Penalty

Modern processors may have  $> 10$  pipeline stages between next PC calculation and branch resolution !

How much work is lost if pipeline doesn't follow correct instruction flow?

$\sim \text{Loop length} \times \text{pipeline width}$





# Average Run-Length between Branches

Average dynamic instruction mix from SPEC92:

	SPECint92	SPECfp92
ALU	39 %	13 %
FPU Add		20 %
FPU Mult		13%
load	26 %	23 %
store	9 %	9 %
branch	16 %	8 %
other	10 %	12 %

SPECint92: *compress, eqntott, espresso, gcc , li*  
SPECfp92: *doduc, ear, hydro2d, mdijdp2, su2cor*





## MIPS Branches and Jumps

Each instruction fetch depends on one or two pieces of information from the preceding instruction:

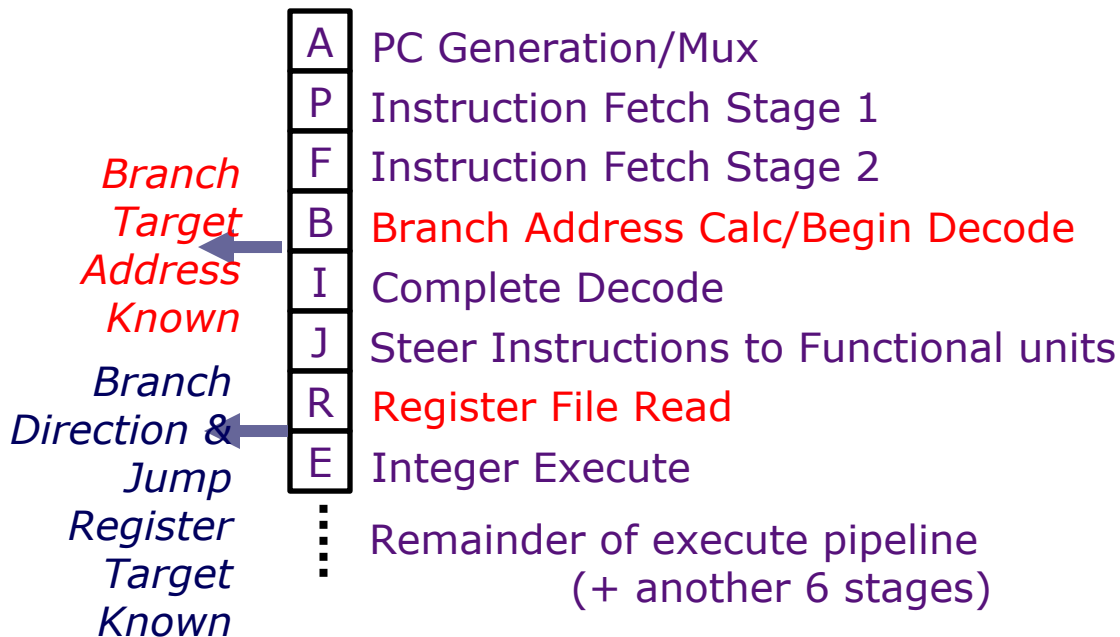
- 1) Is the preceding instruction a **taken branch**?
- 2) If so, what is **the target address**?

J 指令?



# Branch Penalties in Modern Pipelines

UltraSPARC-III instruction fetch pipeline stages  
(in-order issue, 4-way superscalar, 750MHz, 2000)





# Reducing Control Flow Penalty

## Software solutions

- *Eliminate branches* - *loop unrolling*  
Increases the run length
- *Reduce resolution time* - *instruction scheduling*  
Compute the branch condition as early  
as possible (of limited value)

## Hardware solutions

- *Find something else to do* - *delay slots*  
Replaces pipeline bubbles with useful work  
(requires software cooperation)
- *Speculate* - *branch prediction*  
*Speculative execution* of instructions beyond the  
branch



# Branch Prediction

## *Motivation:*

Branch penalties limit performance of deeply pipelined processors

Modern branch predictors have high accuracy (>95%) and can reduce branch penalties significantly

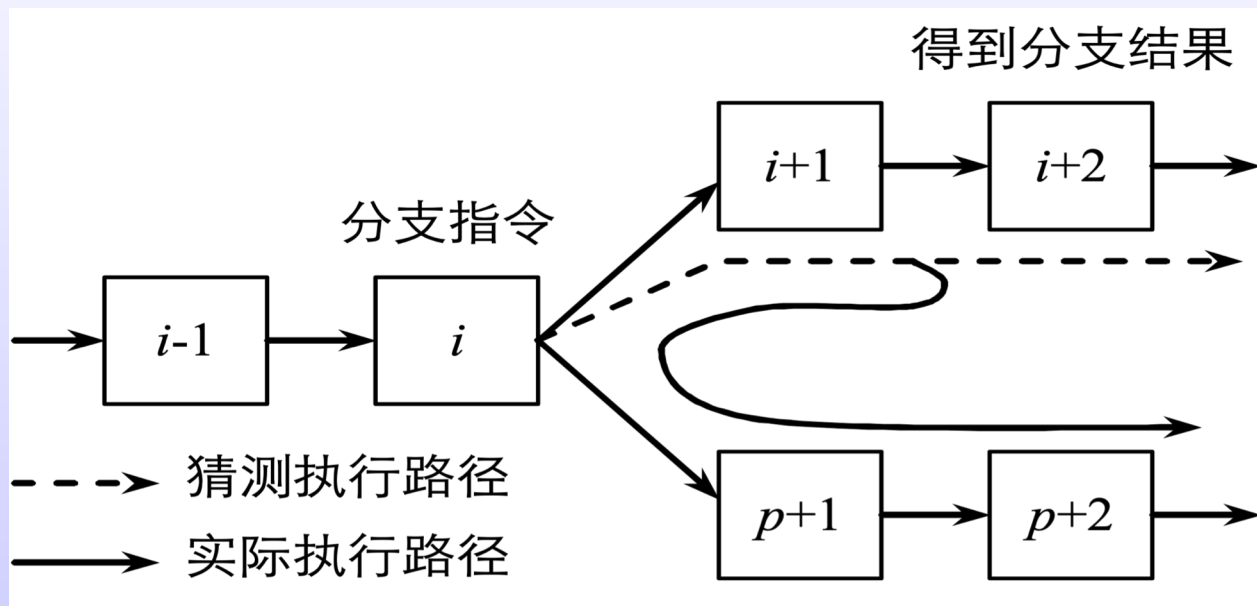
## *Required hardware support:*

### *Prediction structures:*

- Branch history tables, branch target buffers, etc.

### *Mispredict recovery mechanisms:*


- Keep result computation separate from commit
- Kill instructions following branch in pipeline
- Restore state to state following branch





# Static Branch Prediction

Overall probability a branch is taken is ~60-70% but:



ISA can attach preferred direction semantics to branches, e.g., Motorola MC88110


*bne0 (preferred taken) beq0 (not taken)*

ISA can allow arbitrary choice of statically predicted direction, e.g., HP PA-RISC, Intel IA-64

typically reported as ~80% accurate



# Dynamic Prediction



Prediction as a feedback control process

## Operations

- Predict
- Update




# Predictor Primitive

- **Indexed table** holding values

- **Operations**

**Predict**

**Update**




$$\text{Prediction} = P[\text{Width}, \text{Depth}](\text{Index}; \text{Update})$$





# One-bit Predictor



$$A_{21064}(PC; T) = P[1, 2K](PC; T)$$

What happens on loop branches?

At best, mispredicts twice for every use of loop



# Dynamic Branch Prediction

*learning based on past behavior*

## *Temporal correlation*

The way a branch resolves may be a good predictor of the way it will resolve at the next execution

## *Spatial correlation*

Several branches may resolve in a highly correlated manner (*a preferred path of execution*)

```
if (x[i] < 7) then  
    y += 1;  
if (x[i] < 5) then  
    c -= 4;
```



# Branch Prediction

- **Easiest (static prediction)**
  - Always taken, always not taken
  - Opcode based
  - Displacement based (**forward not taken, backward taken**)
  - Compiler directed (branch likely, branch not likely)
- **Next easiest**
  - 1 bit predictor** – remember last taken/not taken per branch
    - Use a **branch-prediction buffer or branch-history table**
    - Use part of the PC (low-order bits) to index buffer/table
      - Multiple branches may share the same bit! ! ! ! !
    - Invert the bit if the prediction is wrong
    - **Backward branches for loops will be mispredicted twice(?)**



# Branch Prediction

**Q:** Assume a loop branch is taken nine times in a row, then not taken once. What is the prediction accuracy using 1-bit predictor?

**A:** After first loop, the predictor will say not to take because the last time the execution came out of loop, it set a “0” in the predictor. So, it’s a misprediction. The bit will now be set to “1”. Works fine until the last loop when it is predicted as taken. So, 2 mispredictions in 10 loop executions => 80% accuracy.

**Actual branches**

ttt ttt ttt n

**Predictions**

011 111 111 1

**How about a 2-bit predictor? Let the prediction be changed only after it misses twice in a row.**



# Branch Prediction Bits


- Assume 2 BP bits per instruction
- Use **saturating counter** or others

On $\neg$ taken $\rightarrow$	$\rightarrow$ On taken	1	1	Strongly taken
		1	0	Weakly taken
		0	1	Weakly $\neg$ taken
		0	0	Strongly $\neg$ taken




# Bimodal Prediction

PC




- **Table of 2-bit saturating counters**  
**Predict the most common direction**

**Advantages: simple, cheap, “good” accuracy**





# Two-bit Predictor (**adder**)

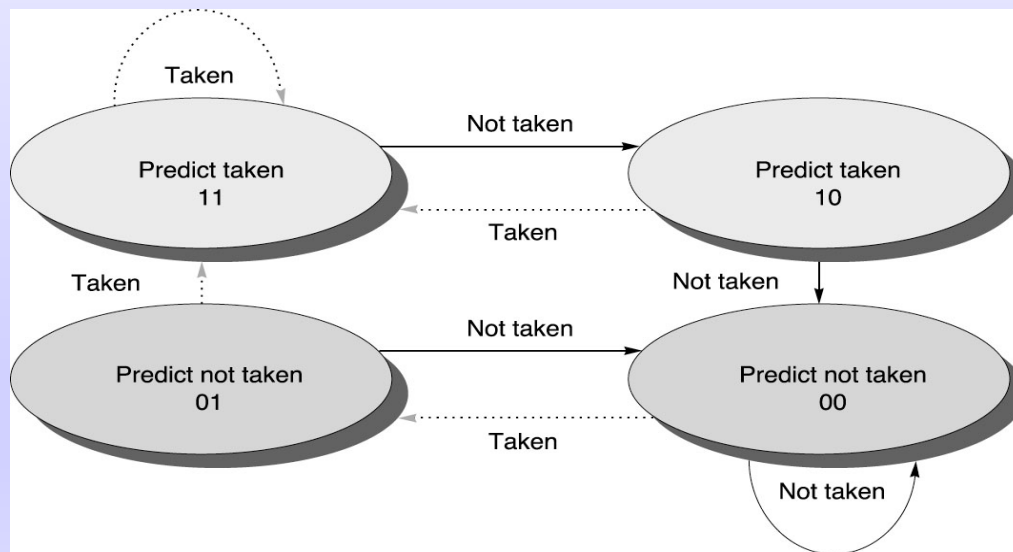


$$\text{Counter}[W,D](I; T) = P[W, D](I; \text{if } T \text{ then } P+1 \text{ else } P-1)$$



## 2-bit Branch Prediction(实现? )

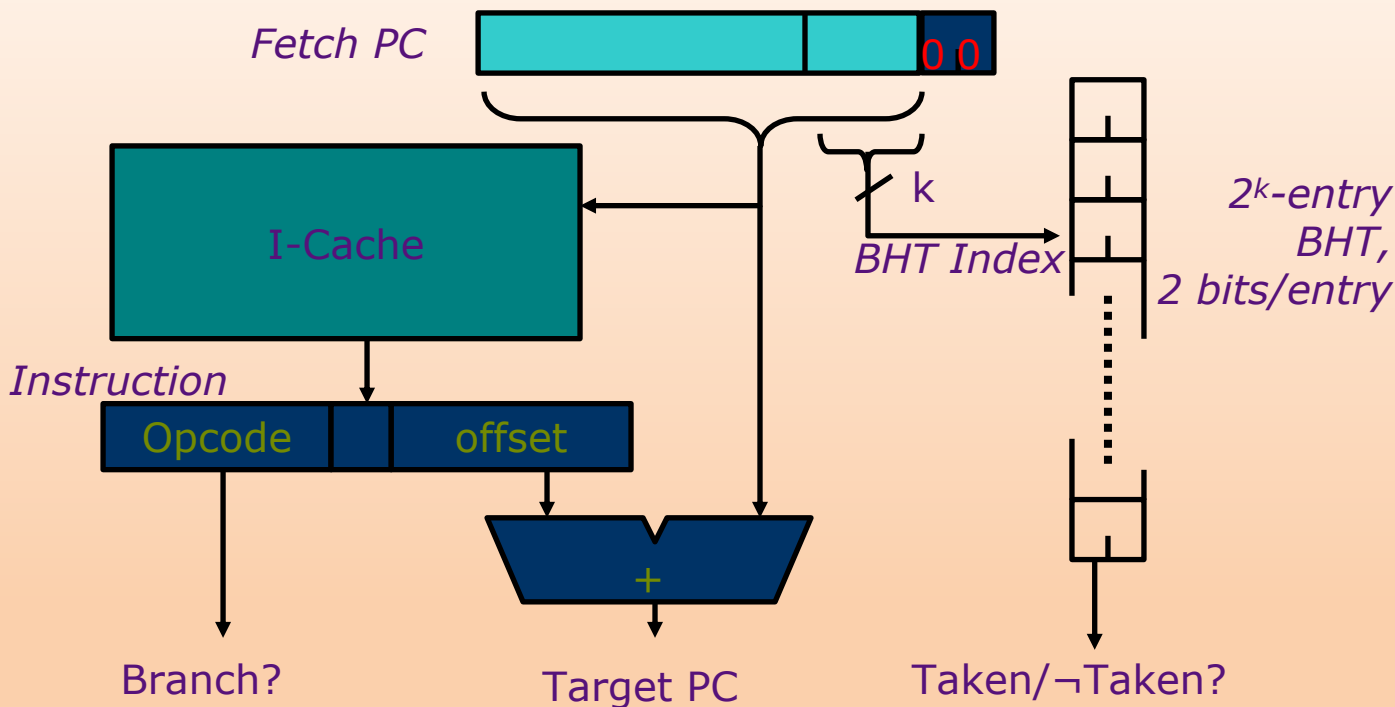
- Has 4 states instead of 2, allowing for more information about tendencies
- A prediction must miss twice before it is changed







# Branch History Table



4K-entry BHT, 2 bits/entry, ~80-90% correct predictions



# 2-bit Branch Prediction

- Good for backward branches of loops

Actual branches


ttt ttt ttt n

Predictions

t t t t t t t t t t

Prediction status bits

10 11 11 11 11 11 11 11 11 11

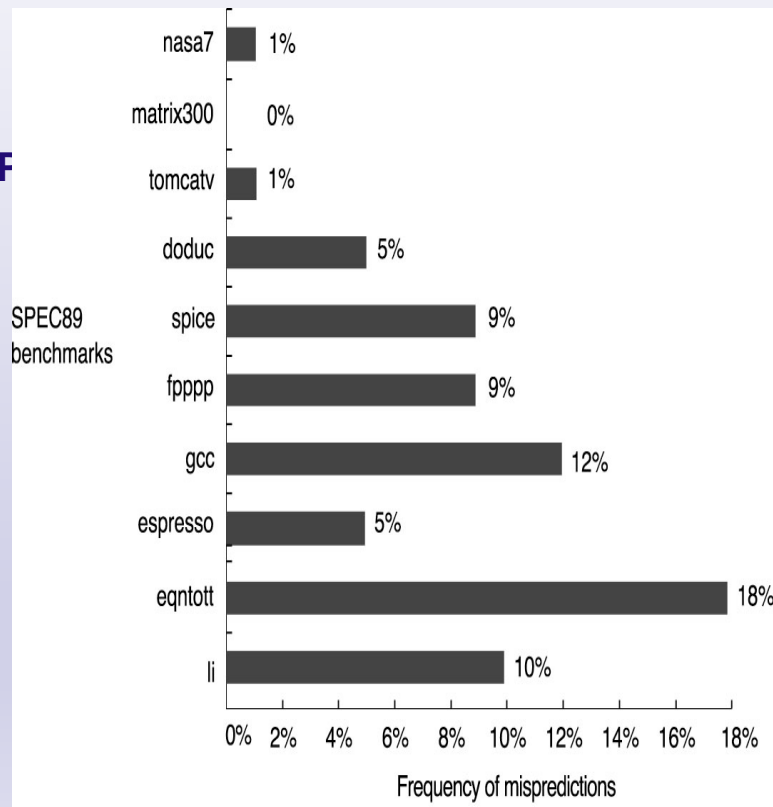



实现？



# Branch History Table

- Has limited size
- 2 bits by N (e.g. 4K)
- Uses low-order bits of branch PC to choose entry





# Branch Target Buffer (BTB)

## Is the current instruction a branch ?


- BTB provides the answer before the current instruction is decoded and therefore enables fetching to begin after **IF-stage** .

## What is the branch target ?

- BTB provides the **branch target** if the prediction is a taken direct branch (for not taken branches the target is simply  $PC+4$  ) .



# Branch Target Buffer

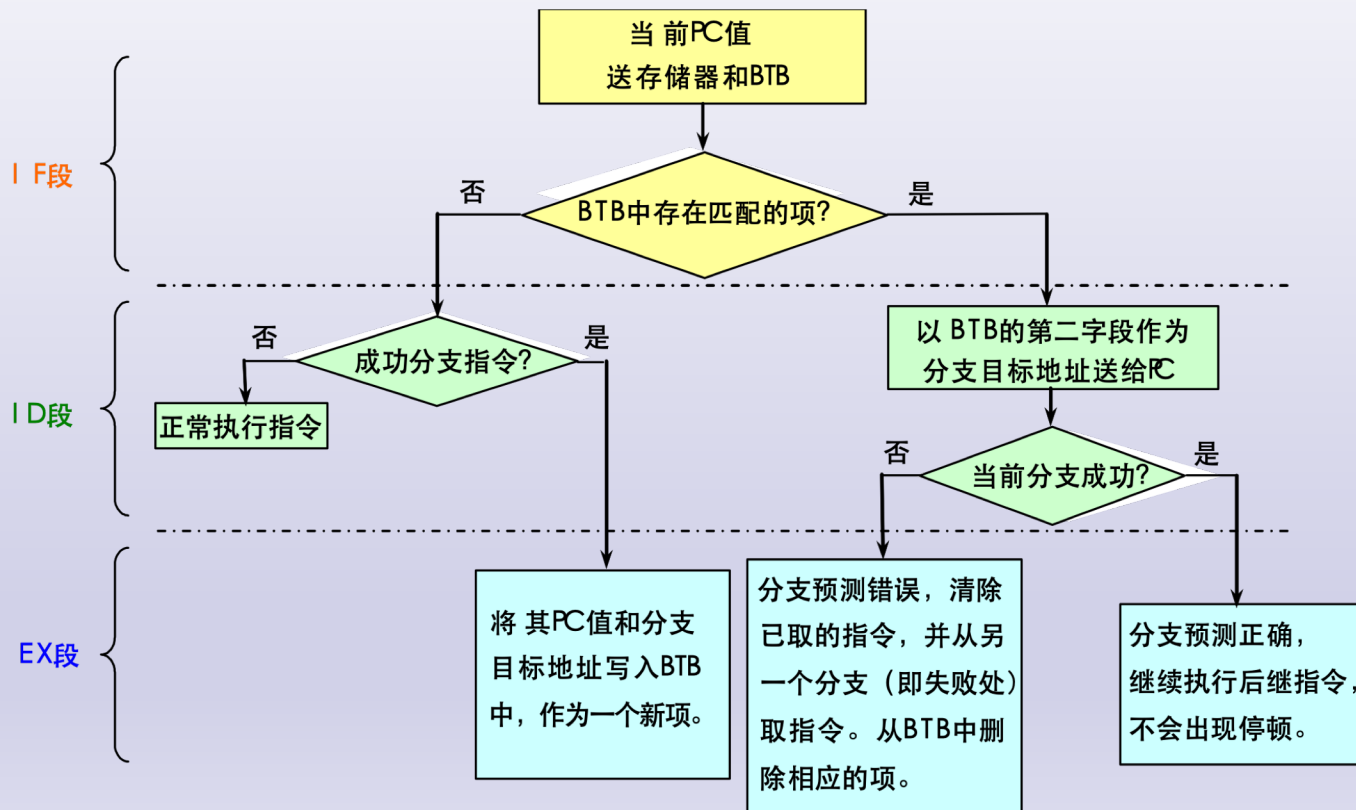


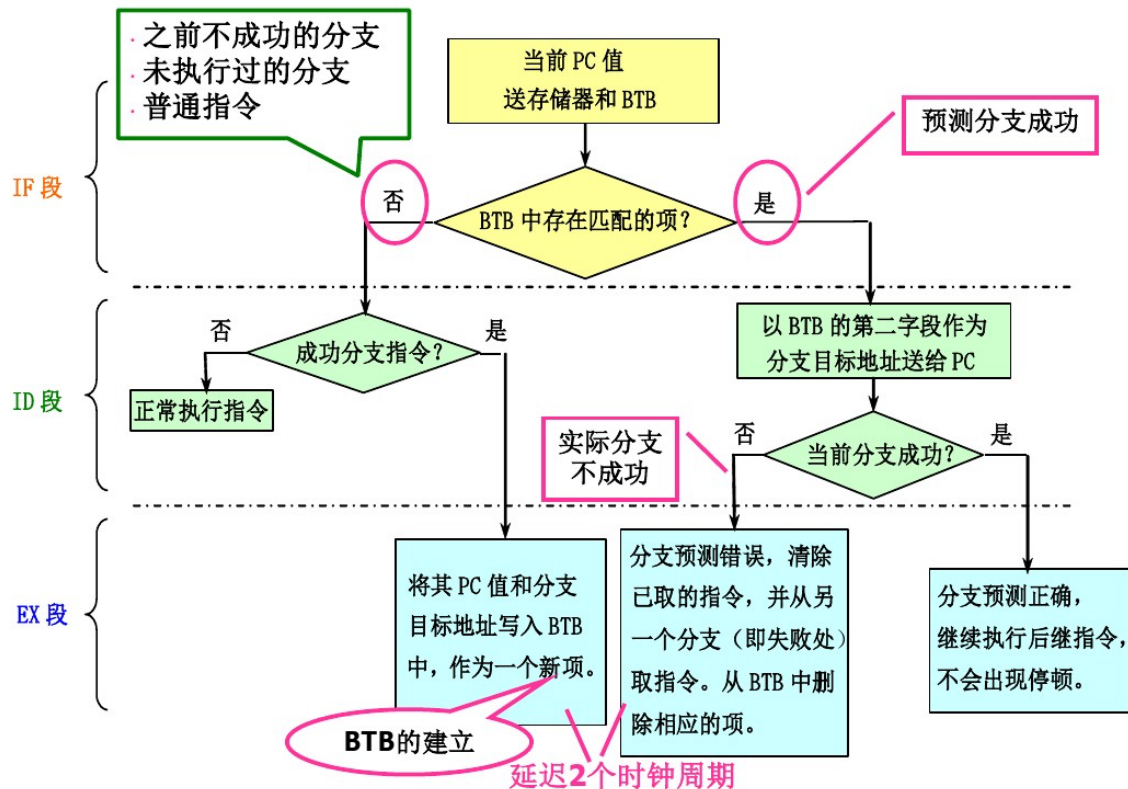
BP bits are stored with the predicted target address.

IF stage: *If (BP=taken) then nPC=target else nPC=PC+4*  
later: *check prediction, if wrong then kill the instruction  
and update BTB & BPb else update BPb*



# Case Study









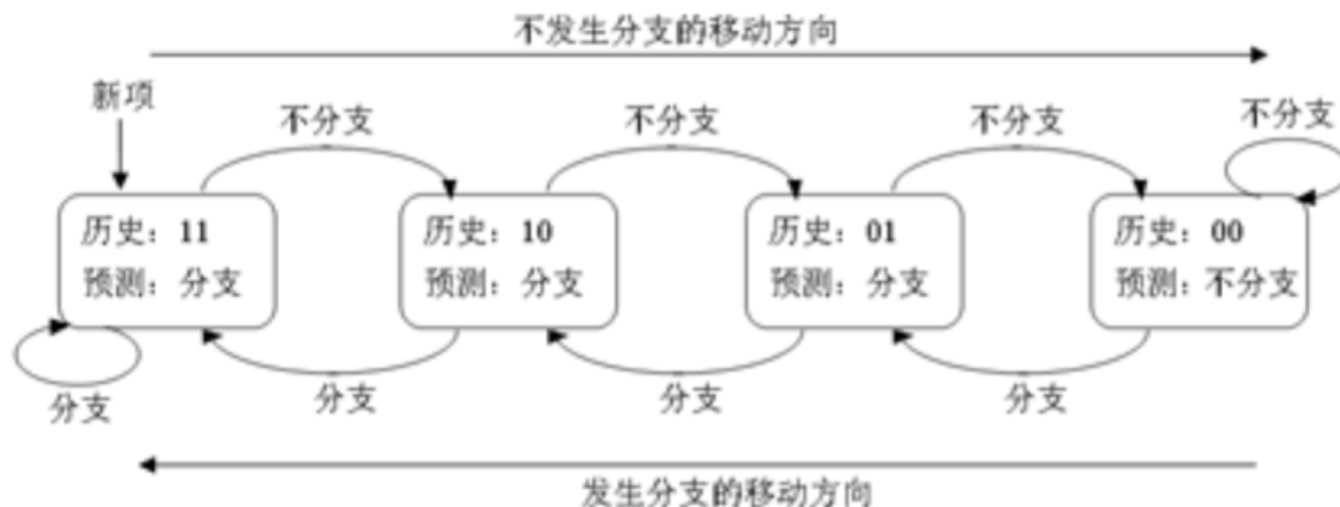
指令在BTB中?	预测	实际情况	延迟周期
是	成功	成功	0
是	成功	不成功	2
不是		成功	2
不是		不成功	0



# BTB-Pentium



(a) 分支目标缓冲器 BTB 项



(b) 分支目标缓冲器 BTB 的历史状态流程